

# **CSCI 3753**

# **Operating Systems**

## **Device Management**

**Lecture Notes By**

**Shivakant Mishra**


**Computer Science, CU-Boulder**

**Last Update: 08/30/2016**

# Recap ...

- System Boot
  - What happens when you switch on or reset a computing system
- Protecting OS from applications
  - How can we prevent an application program from corrupting the operating system
- System Call API
  - How is the system call interface that enables application programs to access OS services implemented

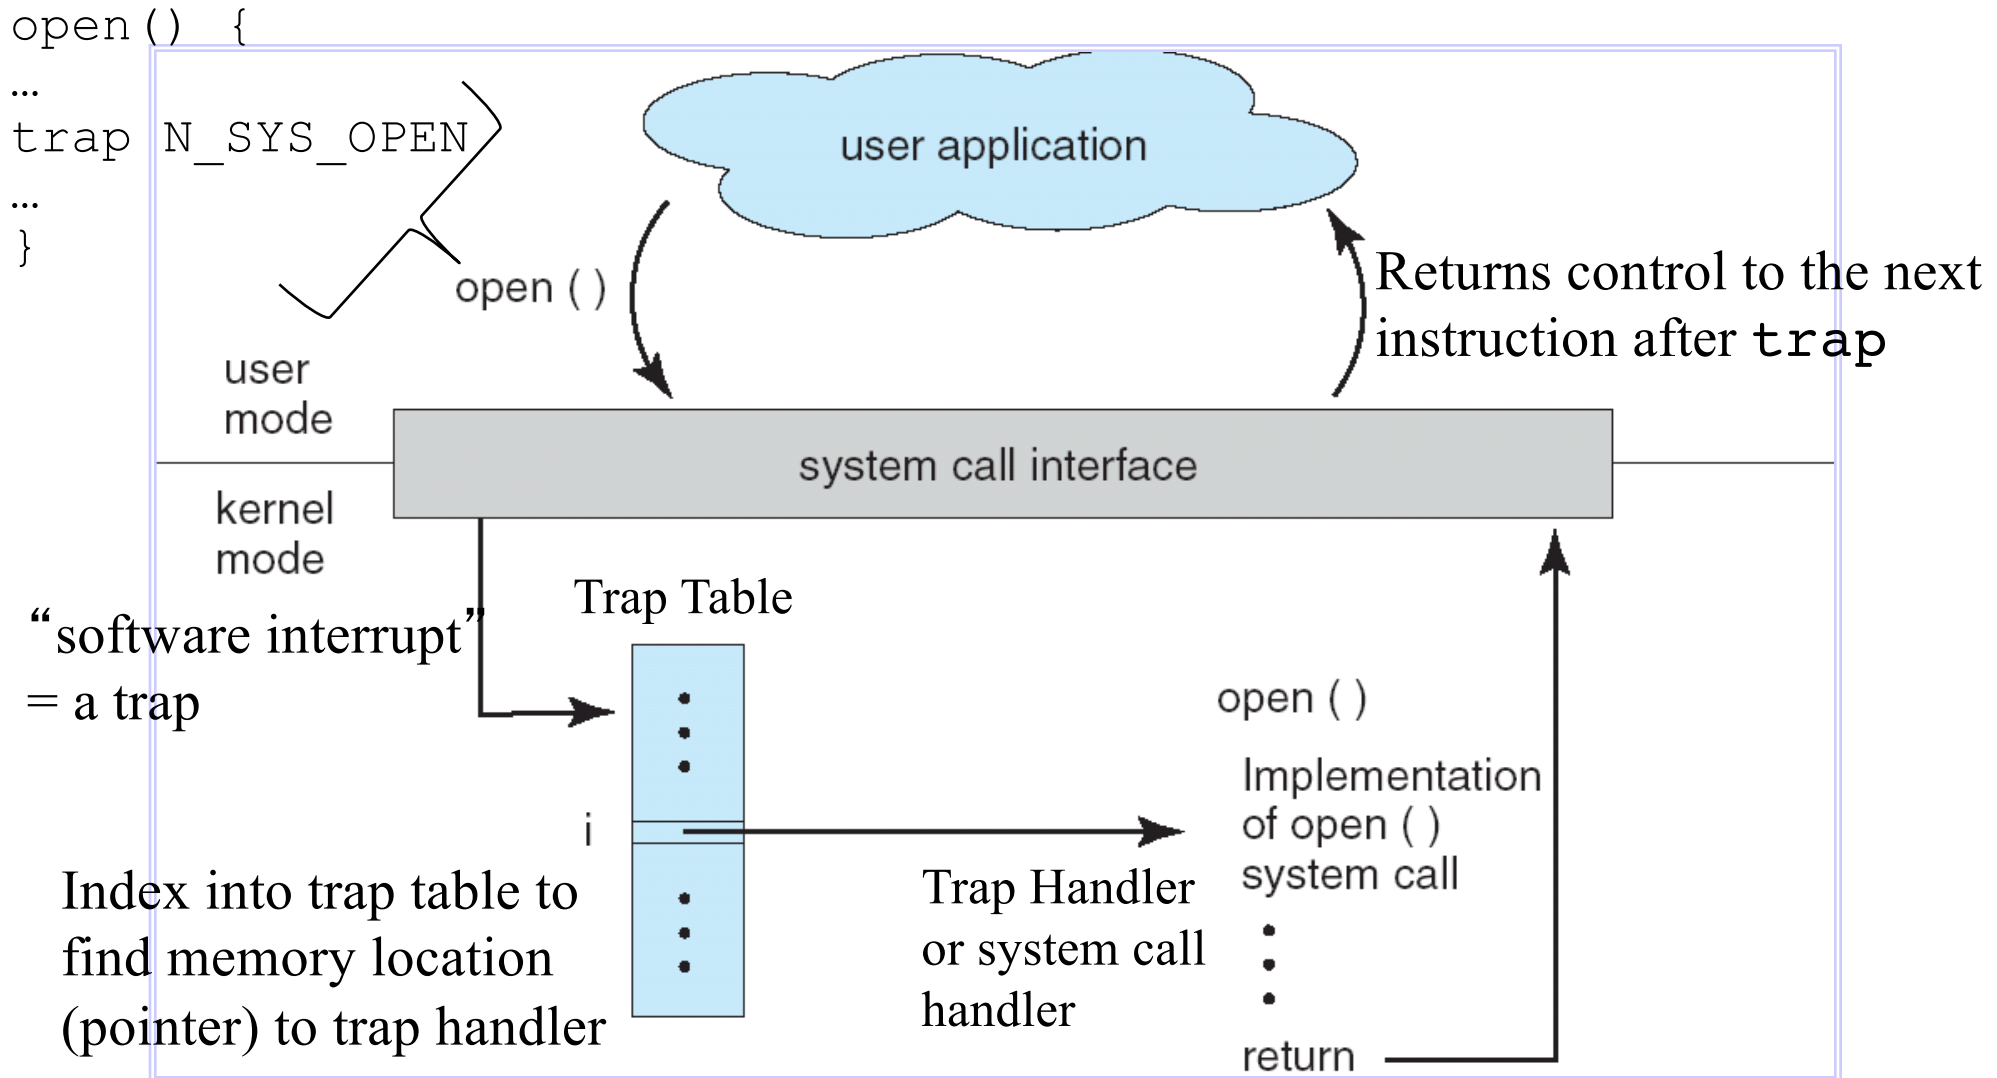
# Intel System Initialization



# Processor mode

- Supervisor mode or user mode:
  - Supervisor mode (mode bit = 0): processor can execute every instruction available in the instruction set.
  - User mode (mode bit = 1): processor can execute only a subset of instructions available in the instruction set.
- Privileged (protected) instructions:
  - Instructions that can be executed only in supervisor mode.
  - I/O instructions
  - Protection and security: privileged load and store instructions
- Used to define two classes of memory space: user space and system space.

# API – System Call – OS Relationship




# Announcements

- Programming assignment one will be posted on Moodle by Wednesday, 08/31
  - Adding a system call; compiling OS kernel
  - You will have one week to work on the assignment
- Readings
  - Read chapters 1 and 2
  - We will start with chapter 13 today ...

# Device Management

# Von Neumann Computer Architecture

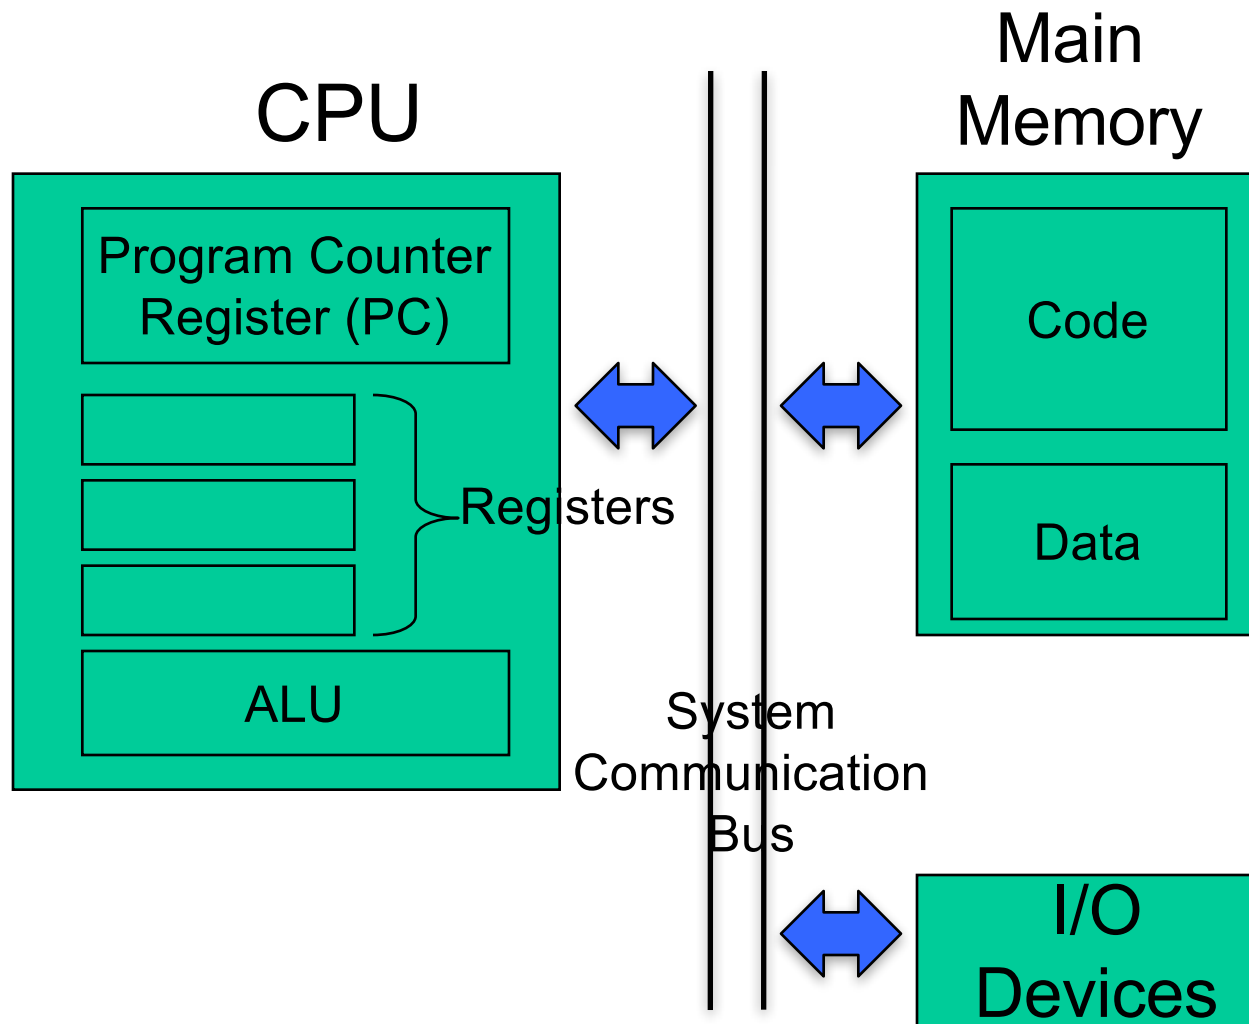


In 1945, von Neumann described a “stored-program” digital computer in which memory stored both instructions \*and\* data

This simplified loading of new programs and executing them without having to rewire the entire computer each time a new program needed to be loaded



# Von Neumann Computer Architecture

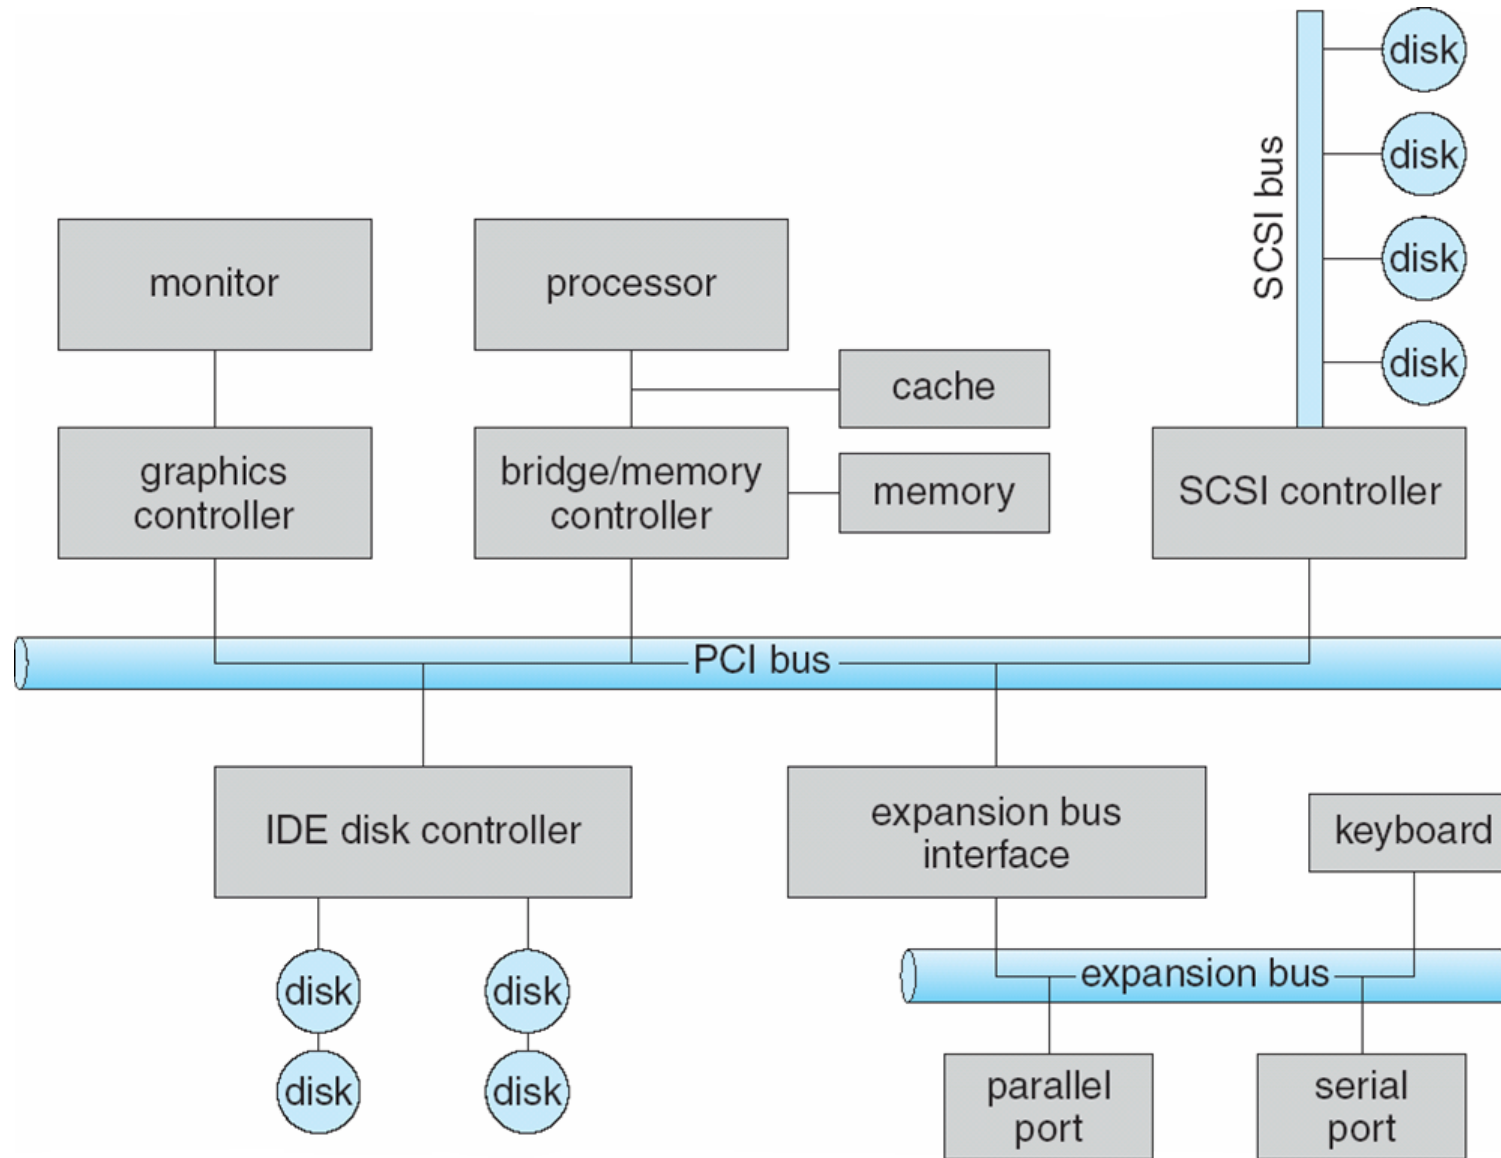


Want to support more devices: card reader, magnetic tape reader, printer, display, disk storage, etc.

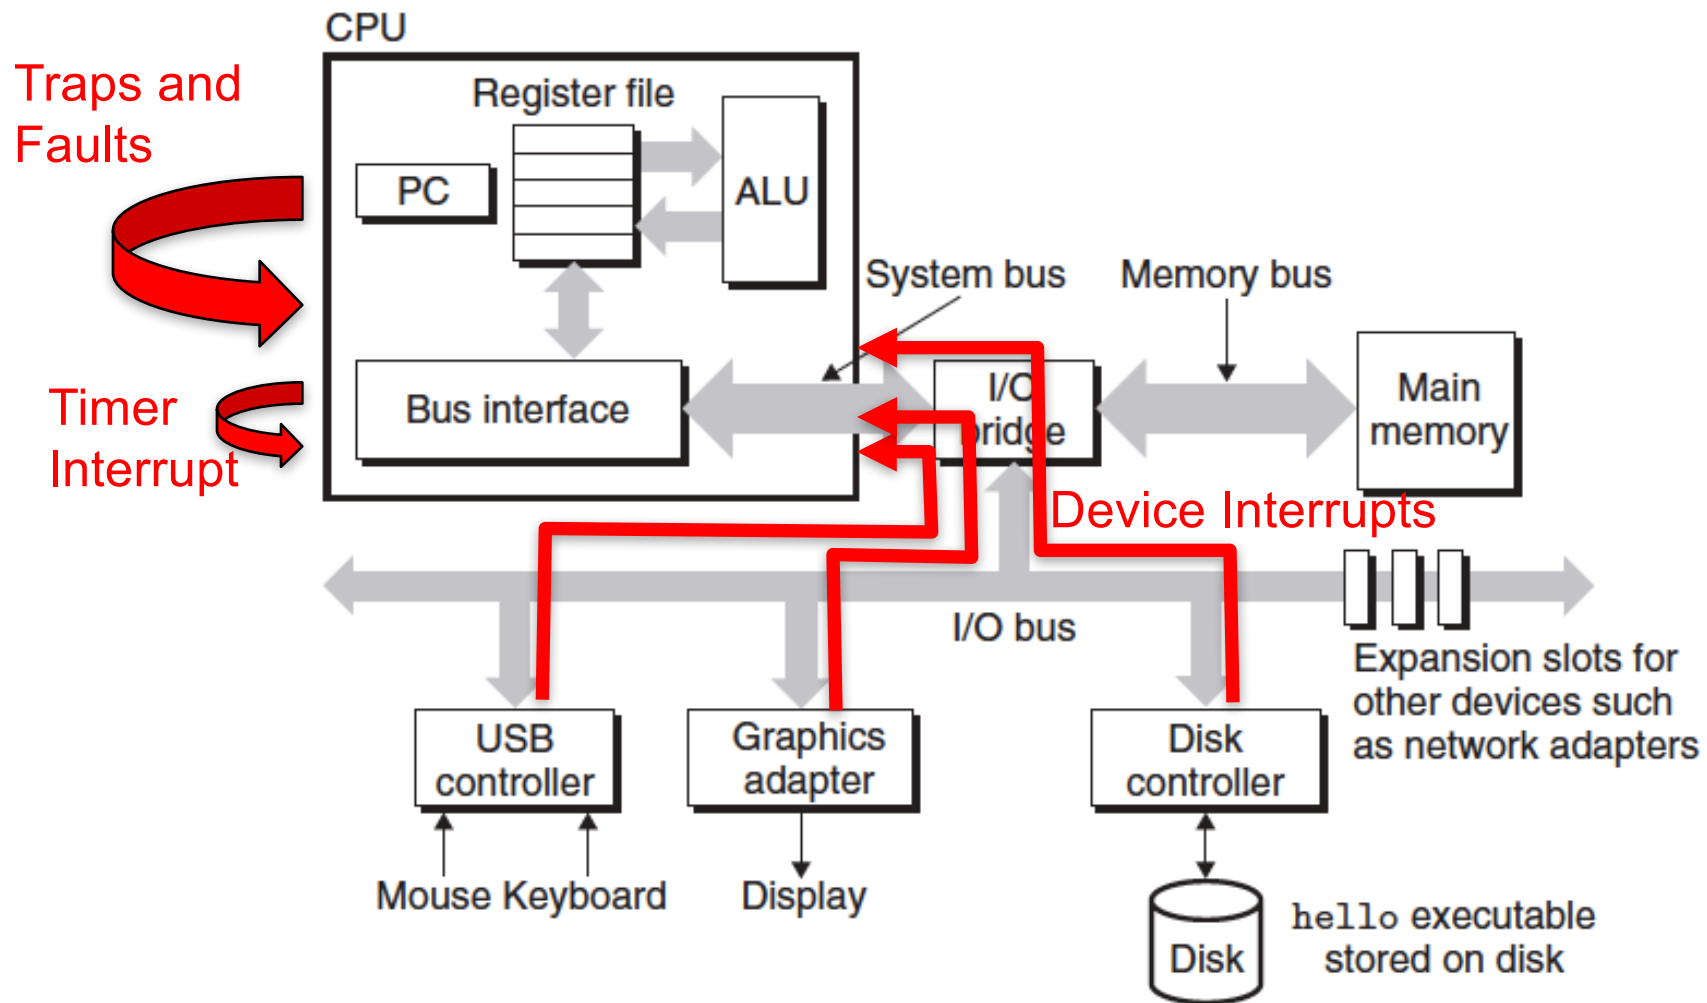
System bus evolved to handle multiple I/O devices.

Includes control, address and data buses

# A Typical PC Bus Structure



# Modern Computer Architecture: Devices and the I/O Bus



# Classes of Exceptions

Class	Cause	Examples	Return behavior
Trap	Intentional exception, i.e. “software interrupt”	System calls	always returns to next instruction, synchronous
Fault	Potentially recoverable error	Divide by 0, stack overflow, invalid opcode, page fault, segmentation fault	might return to current instruction, synchronous
(Hardware) Interrupt	signal from I/O device	Disk read finished, packet arrived on network interface card (NIC)	always returns to next instruction, asynchronous
Abort	nonrecoverable error	Hardware bus failure	never returns, synchronous

# Examples of x86 Exceptions

- x86 Pentium: Table of 256 different exception types
  - some assigned by CPU designers (divide by zero, memory access violations, page faults)
  - some assigned by OS, e.g. interrupts or traps
- Pentium CPU contains exception table base register that points to this table, so it can be located anywhere in memory

# Examples of x86 Exceptions

## Exception Table

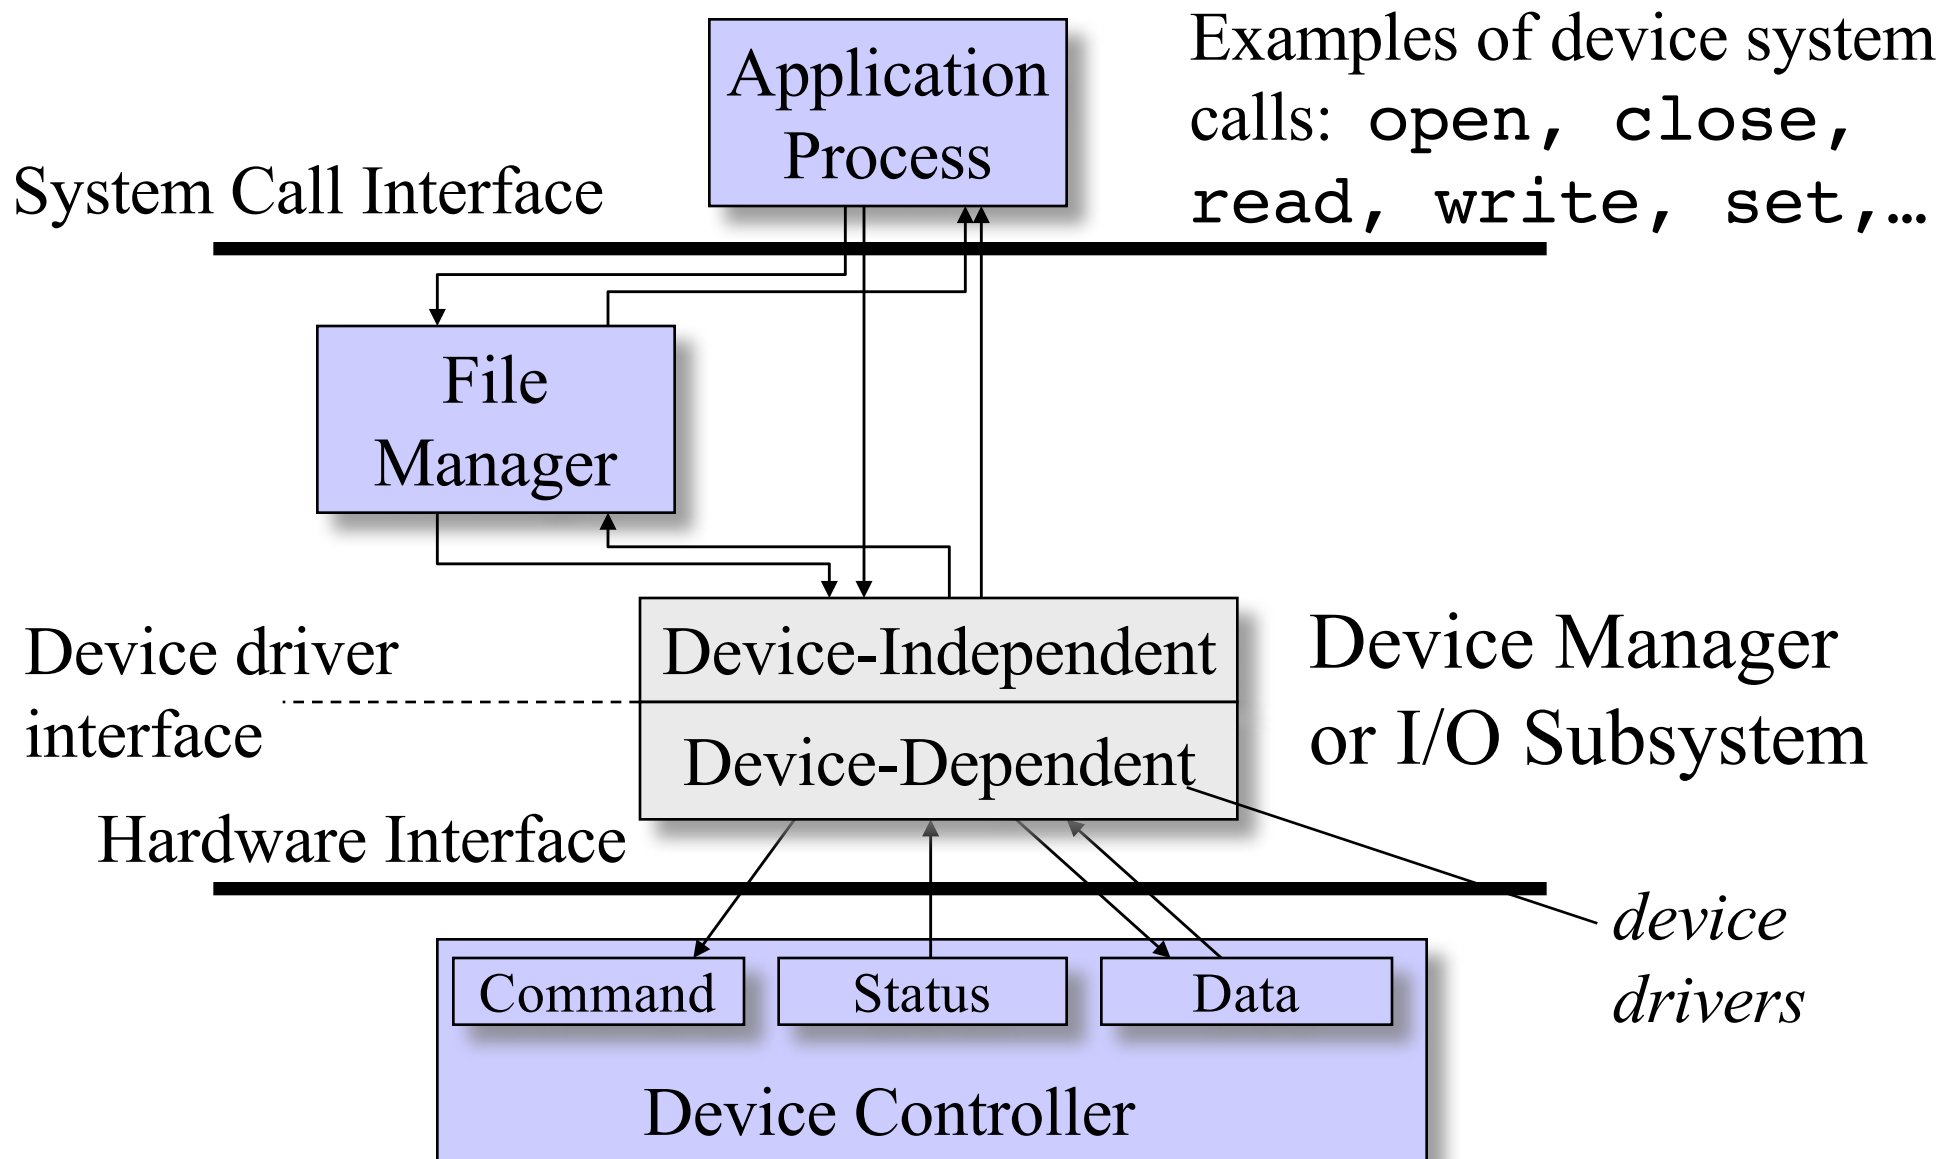
Exception Number	Description	Exception Class	Pointer to Handler
0-31 reserved for hardware	0	Divide error	fault
	13	General protection fault	fault
	14	Page fault	fault
	18	machine check	abort
OS assigns	32-127	OS-defined	Interrupt or trap
	128	System call	Trap
	129-255	OS-defined	Interrupt or trap

offsets form *interrupt vector*

# Device Manager

- Controls operation of I/O devices
  - Issue I/O commands to the devices
  - Catch interrupts
  - Handle errors
  - Provide a simple and easy-to-use interface
    - Device independence: same interface for all devices.

# Device Management Organization



Operating Systems: A Modern Perspective



# Device System Call Interface

- Create a simple standard interface to access most devices
  - Every I/O device driver should support the following: open, close, read, write, set (ioctl in UNIX), stop, etc.
  - Block vs character
  - Sequential vs direct/random access
  - Blocking versus Non-Blocking I/O
    - blocking system call: process put on wait queue until I/O completes
    - non-blocking system call: returns immediately with partial number of bytes transferred, e.g. keyboard, mouse, network
  - Synchronous versus asynchronous
    - asynchronous returns immediately, but at some later time, the full number of bytes requested is transferred

# ioctl and fcntl (input/output control)

- Want a richer interface for managing I/O devices than just open, close, read, write, ...
- ioctl allows a user-space application to configure parameters and/or actions of an I/O device
  - e.g set the speed of a device, or eject a disk
- Usage: *int ioctl(int fd, int cmd, ...)*;
  - Invokes a system call to execute device-specific *cmd* on I/O device *fd*
  - Used for I/O operations and other operations which cannot be expressed by regular system calls
  - Requests are directed to the correct device driver

# ioctl and fcntl (input/output control)

- Avoids having to create new system calls for each new device and/or unforeseen device function
  - Helps make the OS/kernel extensible
- UNIX, Linux, MacOS X all support ioctl, and Windows has its own version
- In UNIX, each device is modeled as a file
  - *fcntl* for file control is related to *ioctl* and is used for configuring file parameters, hence in many cases I/O communication
  - e.g. use *fcntl* to set a network socket to non-blocking
  - part of POSIX API, so portable across platform


# Device Characteristics

- I/O devices consist of two high-level components
  - Mechanical component
  - Electronic component: device controllers
- OS deals with device controllers

# Device Drivers

- Support the device system call interface functions open, read, write, etc. for that device
- Interact directly with the device controllers
  - Know the details of what commands the device can handle, how to set/get bits in device controller registers, etc.
  - Are part of the device-dependent component of the device manager
- Control flow:
  - An I/O system call traps to the kernel, invoking the trap handler for I/O (the device manager), which indexes into a table using the arguments provided to run the correct device driver

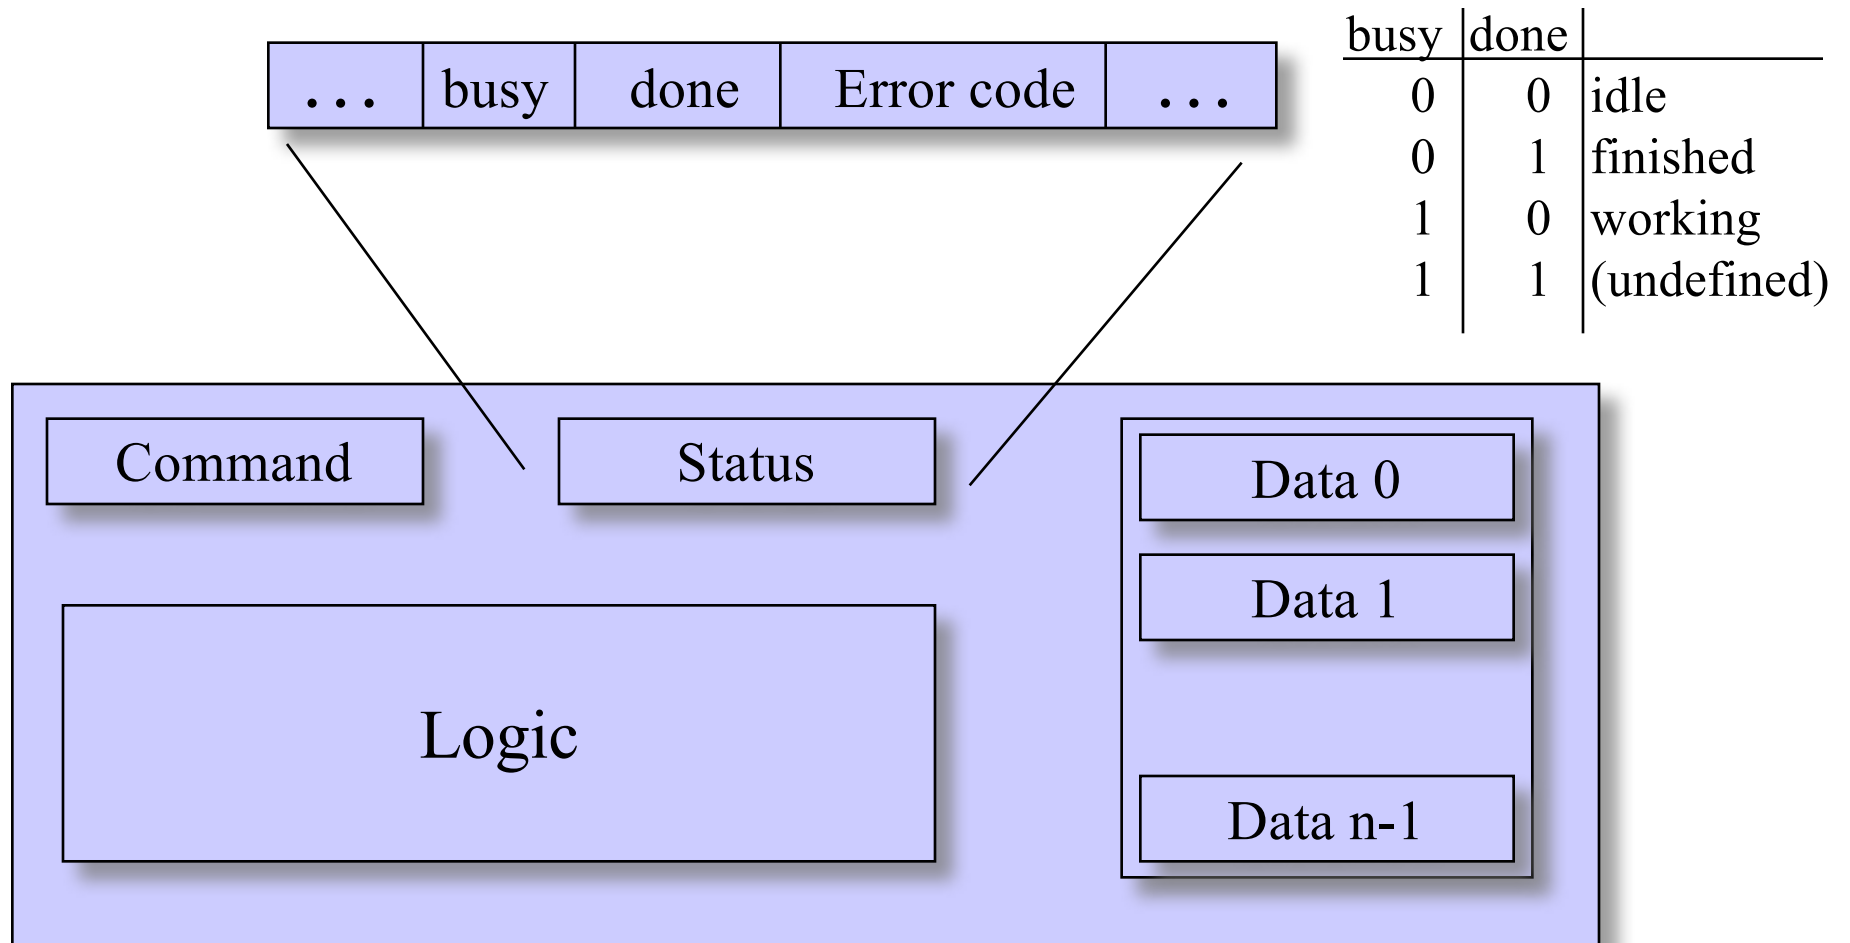
# Device Controller States



- Therefore, need 2 bits for 3 states:
  - A BUSY flag and a DONE flag
  - BUSY=0, DONE=0 => Idle
  - BUSY=1, DONE=0 => Working
  - BUSY=0, DONE=1 => Finished
  - BUSY=1, DONE=1 => Undefined

- Need three states to distinguish the following:
  - Idle: no app is accessing the device
  - Working: one app only is accessing the device
  - Finished: the results are ready for that one app

# Device Controller Interface

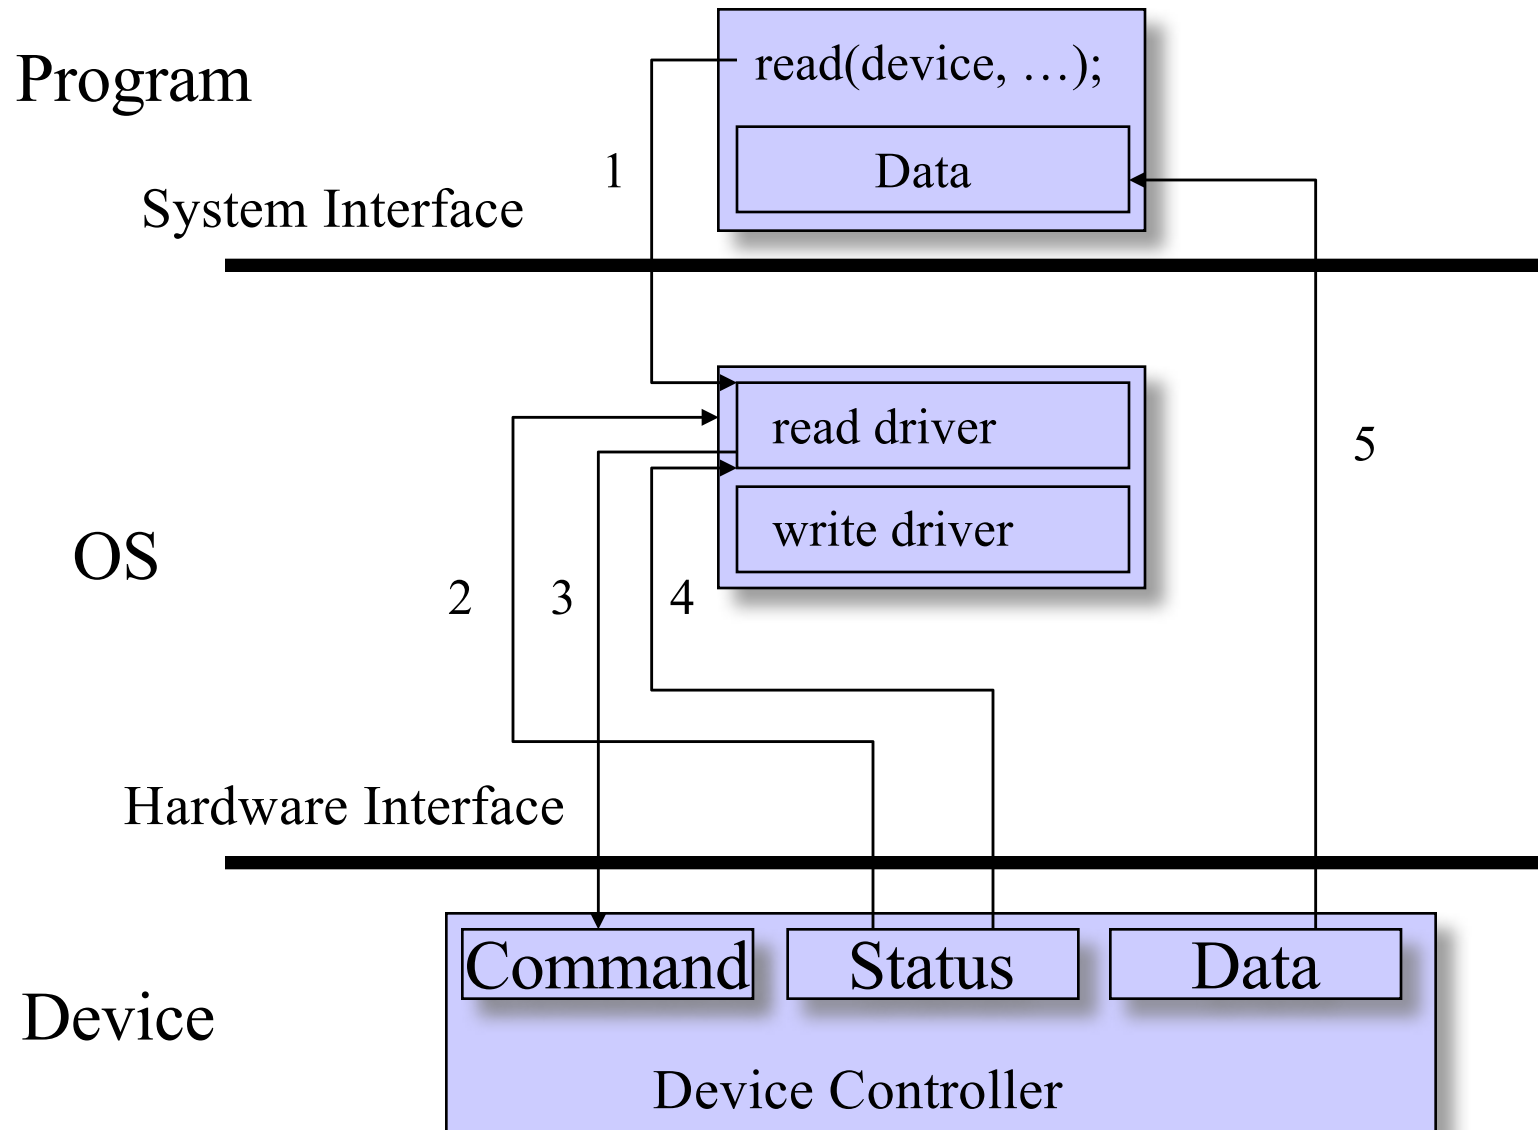


# Polling I/O: A Write Example

	BUSY	DONE
<code>while(deviceN.busy    deviceN.done) &lt;waiting&gt;;</code>	*	*
<code>deviceN.data[0] = &lt;value to write&gt;</code>	0	0
<code>deviceN.command = WRITE;</code>		
<code>while(deviceN.busy) &lt;waiting&gt;;</code>	1	0
<code>/* finished, so read some status bits... */</code>	0	1
<code>deviceN.done = FALSE;</code>	0	0



# Polling I/O Read Operation



# Polling I/O – Busy Waiting

- Note that the OS is spinning in a loop twice:
  - Checking for the device to become idle
  - Checking for the device to finish the I/O request, so the results can be retrieved
  - This wastes CPU cycles that could be devoted to executing applications
- Instead, want to *overlap* CPU and I/O
  - Free up the CPU while the I/O device is processing a read/write

# Device Manager I/O Strategies

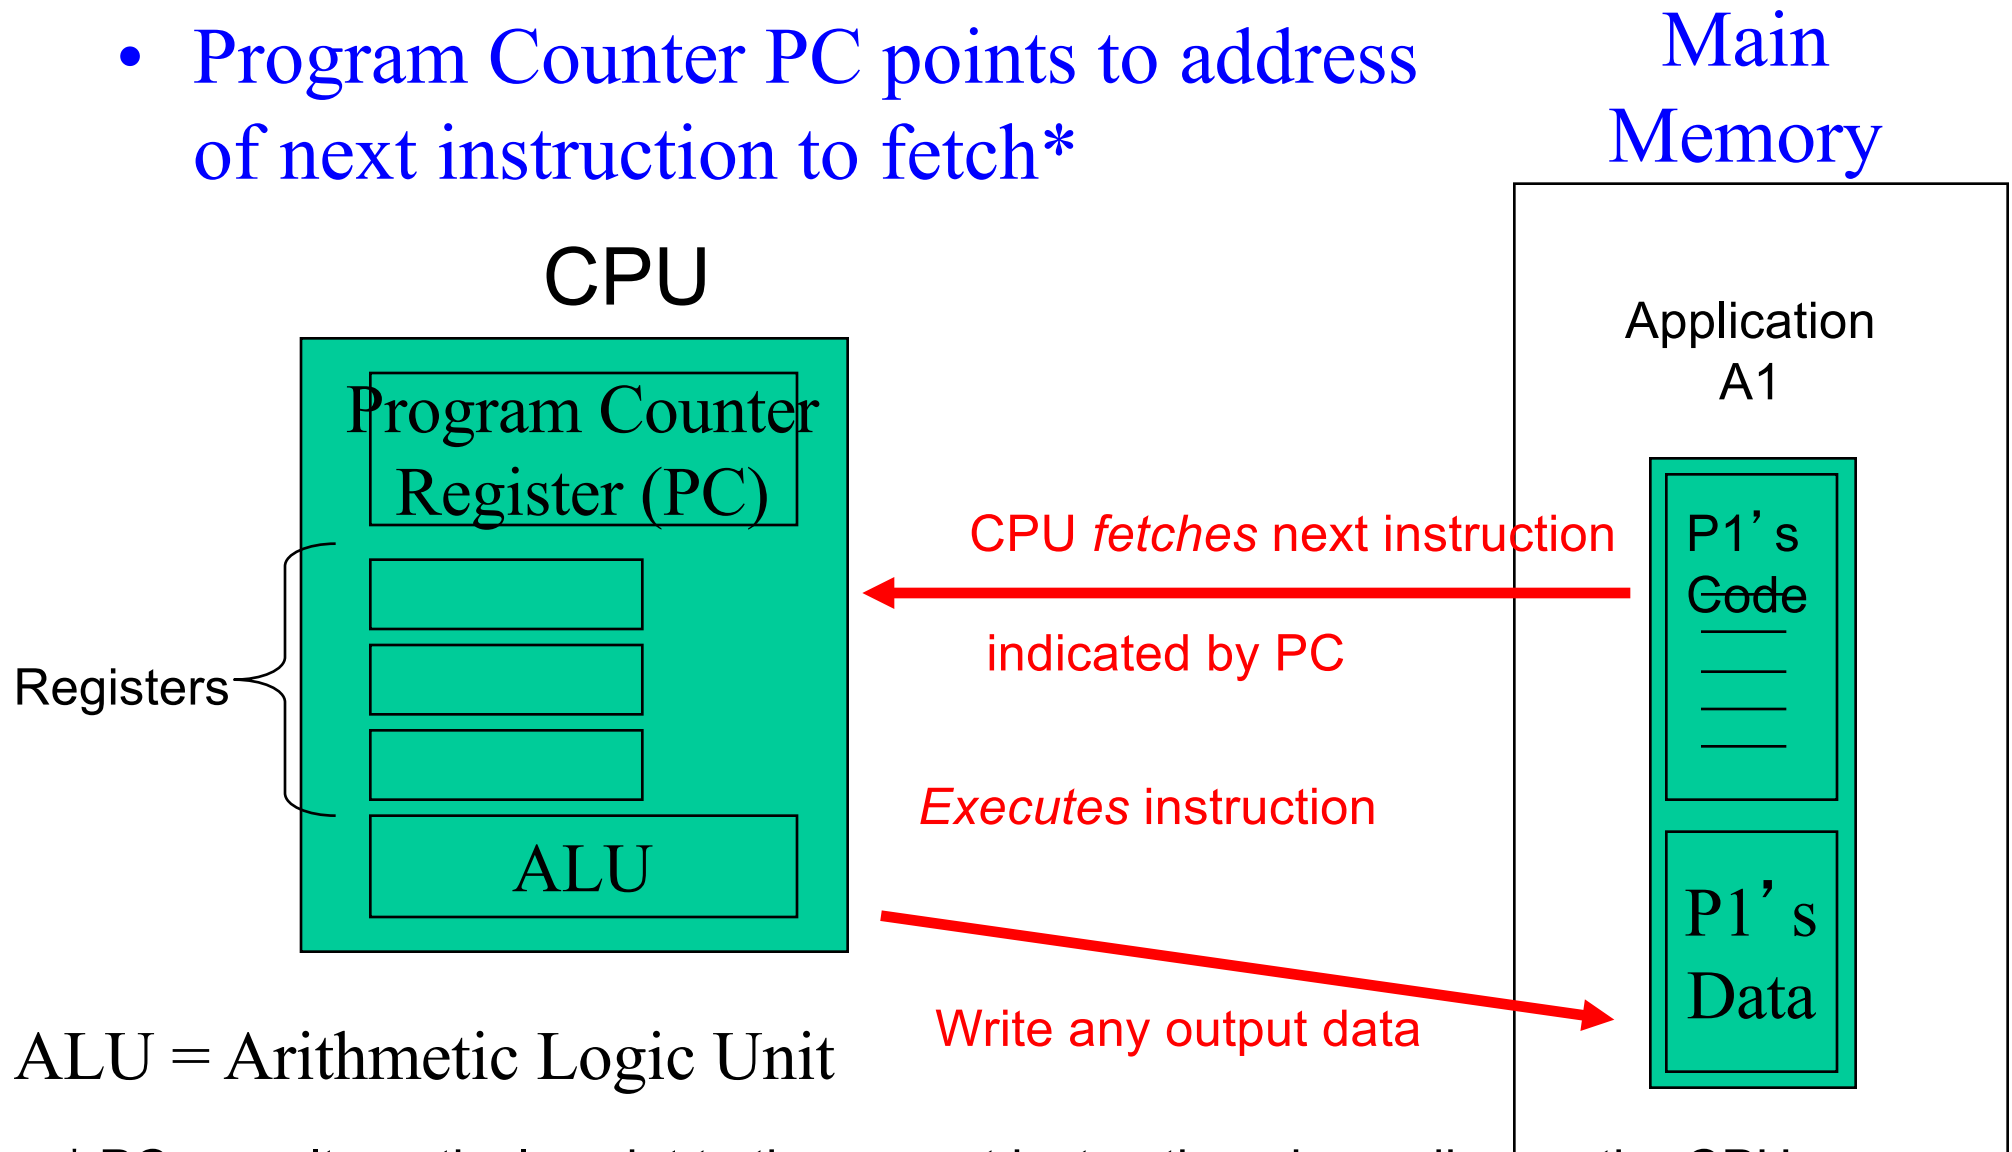
- Underneath the blocking/non-blocking synchronous/asynchronous system call API, OS can implement several strategies for I/O with devices
  - direct I/O with polling
    - the OS device manager busy-waits, we've already seen this
  - direct I/O with *interrupts*
    - More efficient than busy waiting
  - DMA with interrupts

# Hardware Interrupts

- CPU incorporates a *hardware interrupt flag*
- Whenever a device is finished with a read/write, it communicates to the CPU and raises the flag
  - Frees up CPU to execute other tasks without having to keep polling devices
- Upon an interrupt, the CPU interrupts normal execution, and invokes the OS's *interrupt handler*
  - Eventually, after the interrupt is handled and the I/O results processed, the OS resumes normal execution

# CPU Execution of a Program

- Program Counter PC points to address of next instruction to fetch\*



ALU = Arithmetic Logic Unit

\* PC can alternatively point to the current instruction, depending on the CPU

# CPU Checks Interrupt Flag Every Fetch/Execute Cycle

## CPU Pseudocode

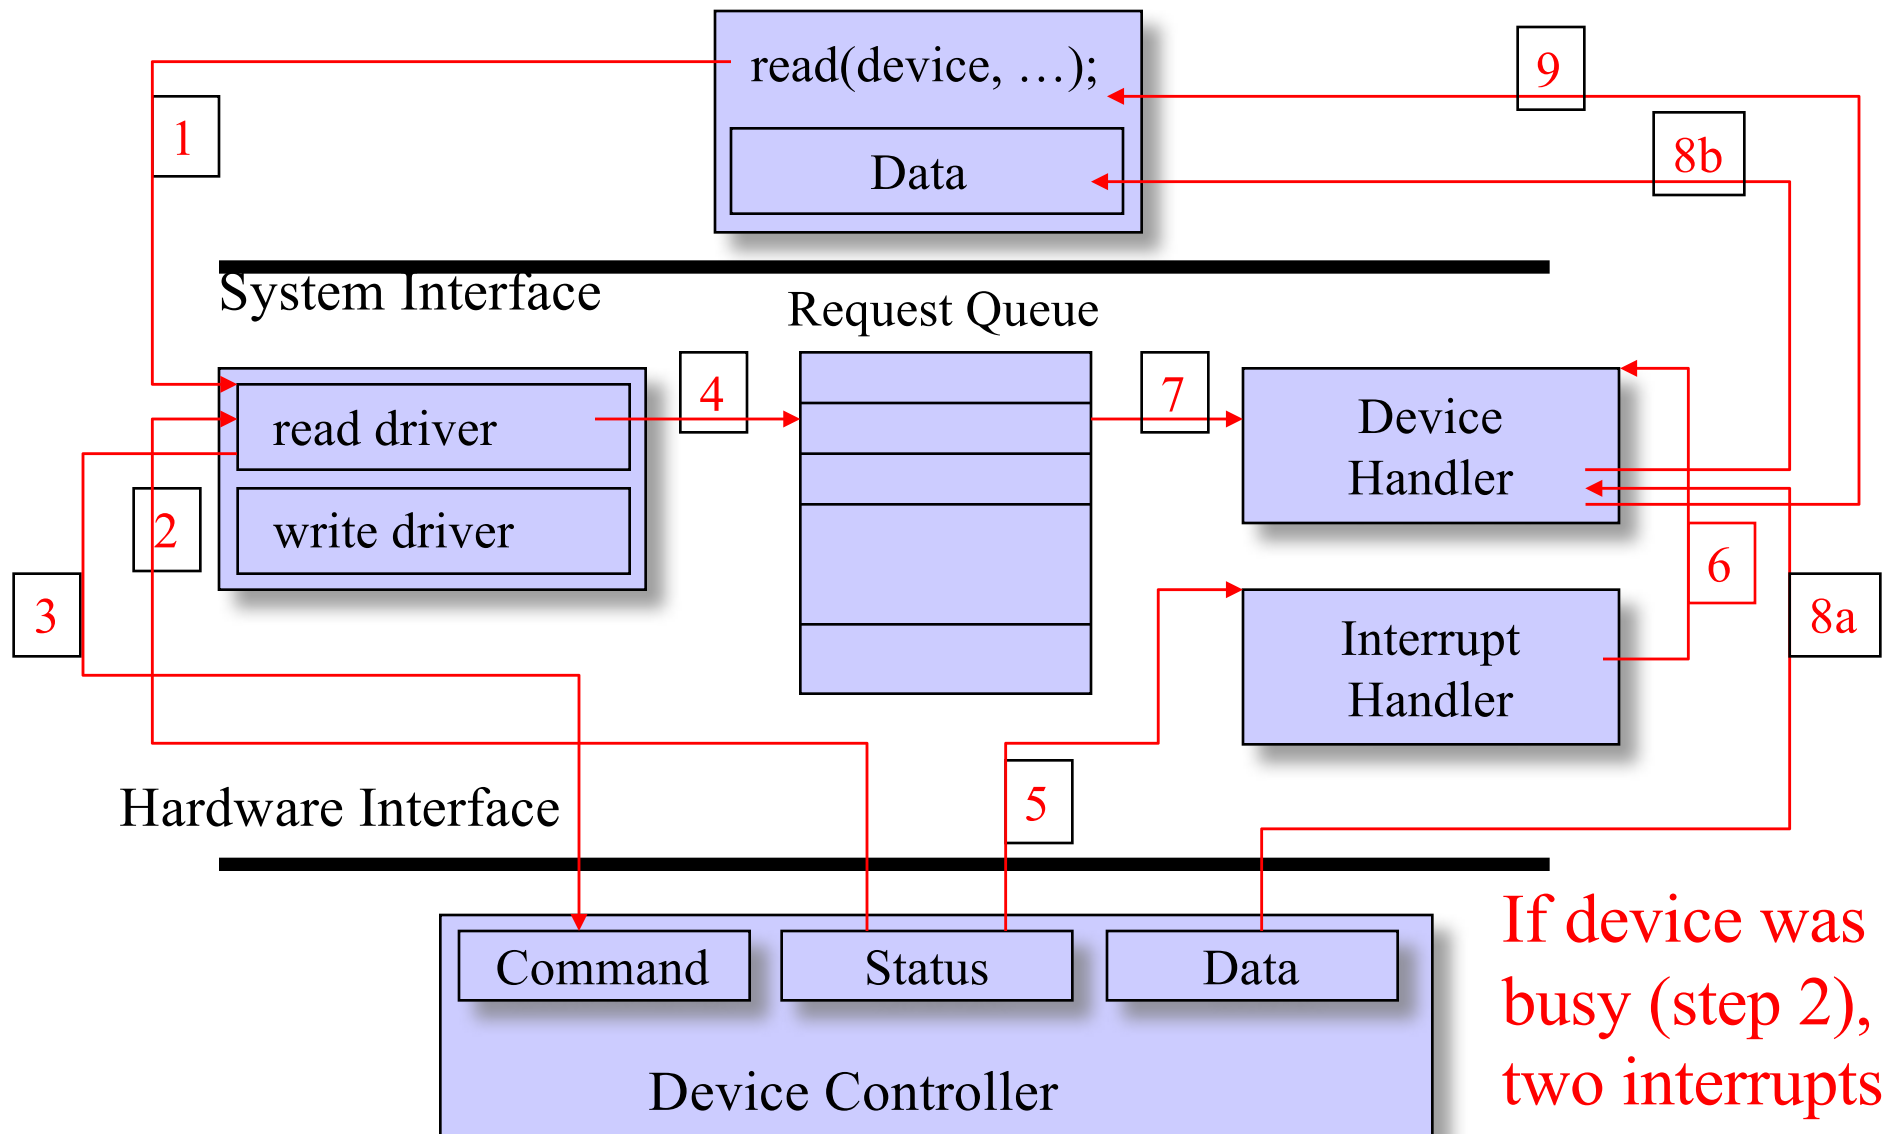
- While (no hardware failure)
  - Fetch next instruction, put in instruction register
  - Execute instruction
  - Check for interrupt: If interrupt flag enabled,
    - Save PC\*
    - Jump to interrupt handler

\* insight from Nutt's text

# Interrupt Handler

- First, save the processor state
  - Save the executing app's program counter (PC) and CPU register data
- Next, find the device causing the interrupt
  - Consult interrupt controller to find the interrupt offset, or poll the devices
- Then, jump to the appropriate device handler
  - Index into the Interrupt Vector using the interrupt offset
  - An Interrupt Service Routine (ISR) either refers to the interrupt handler, or the device handler
- Finally, reenables interrupts

# Interrupt-Driven I/O Operation

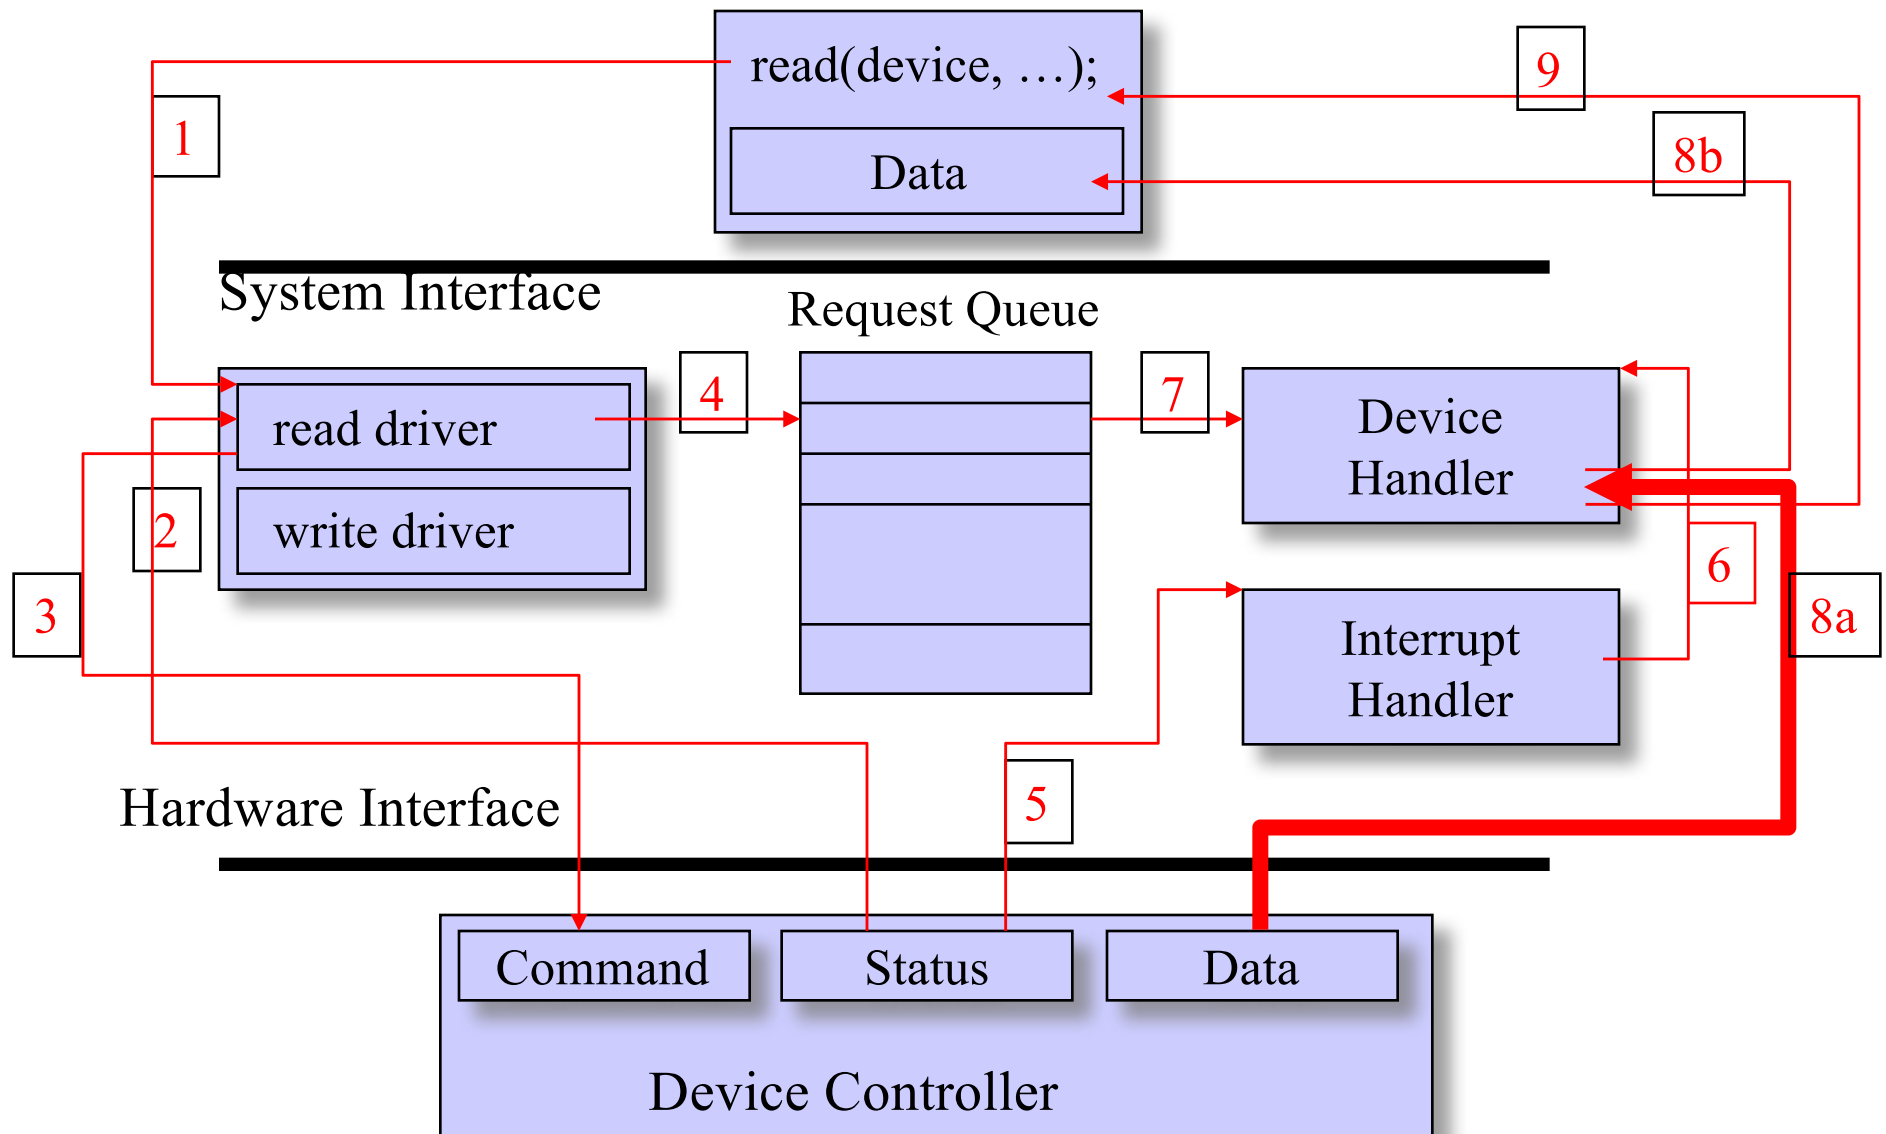




# Problem with Interrupt driven I/O

- Data transfer from disk can become a bottleneck if there is a lot of I/O copying data back and forth between memory and devices
  - Example: want to copy a 1 MB file from disk into memory. The disk is only capable of delivering memory in say 1 KB blocks. So every time a 1 KB block is ready to be copied, an interrupt is raised, interrupting the CPU. This will slow down execution of normal programs and the OS.
  - Worst cases: CPU could be interrupted after the transfer of every byte/character, or every packet from the network card

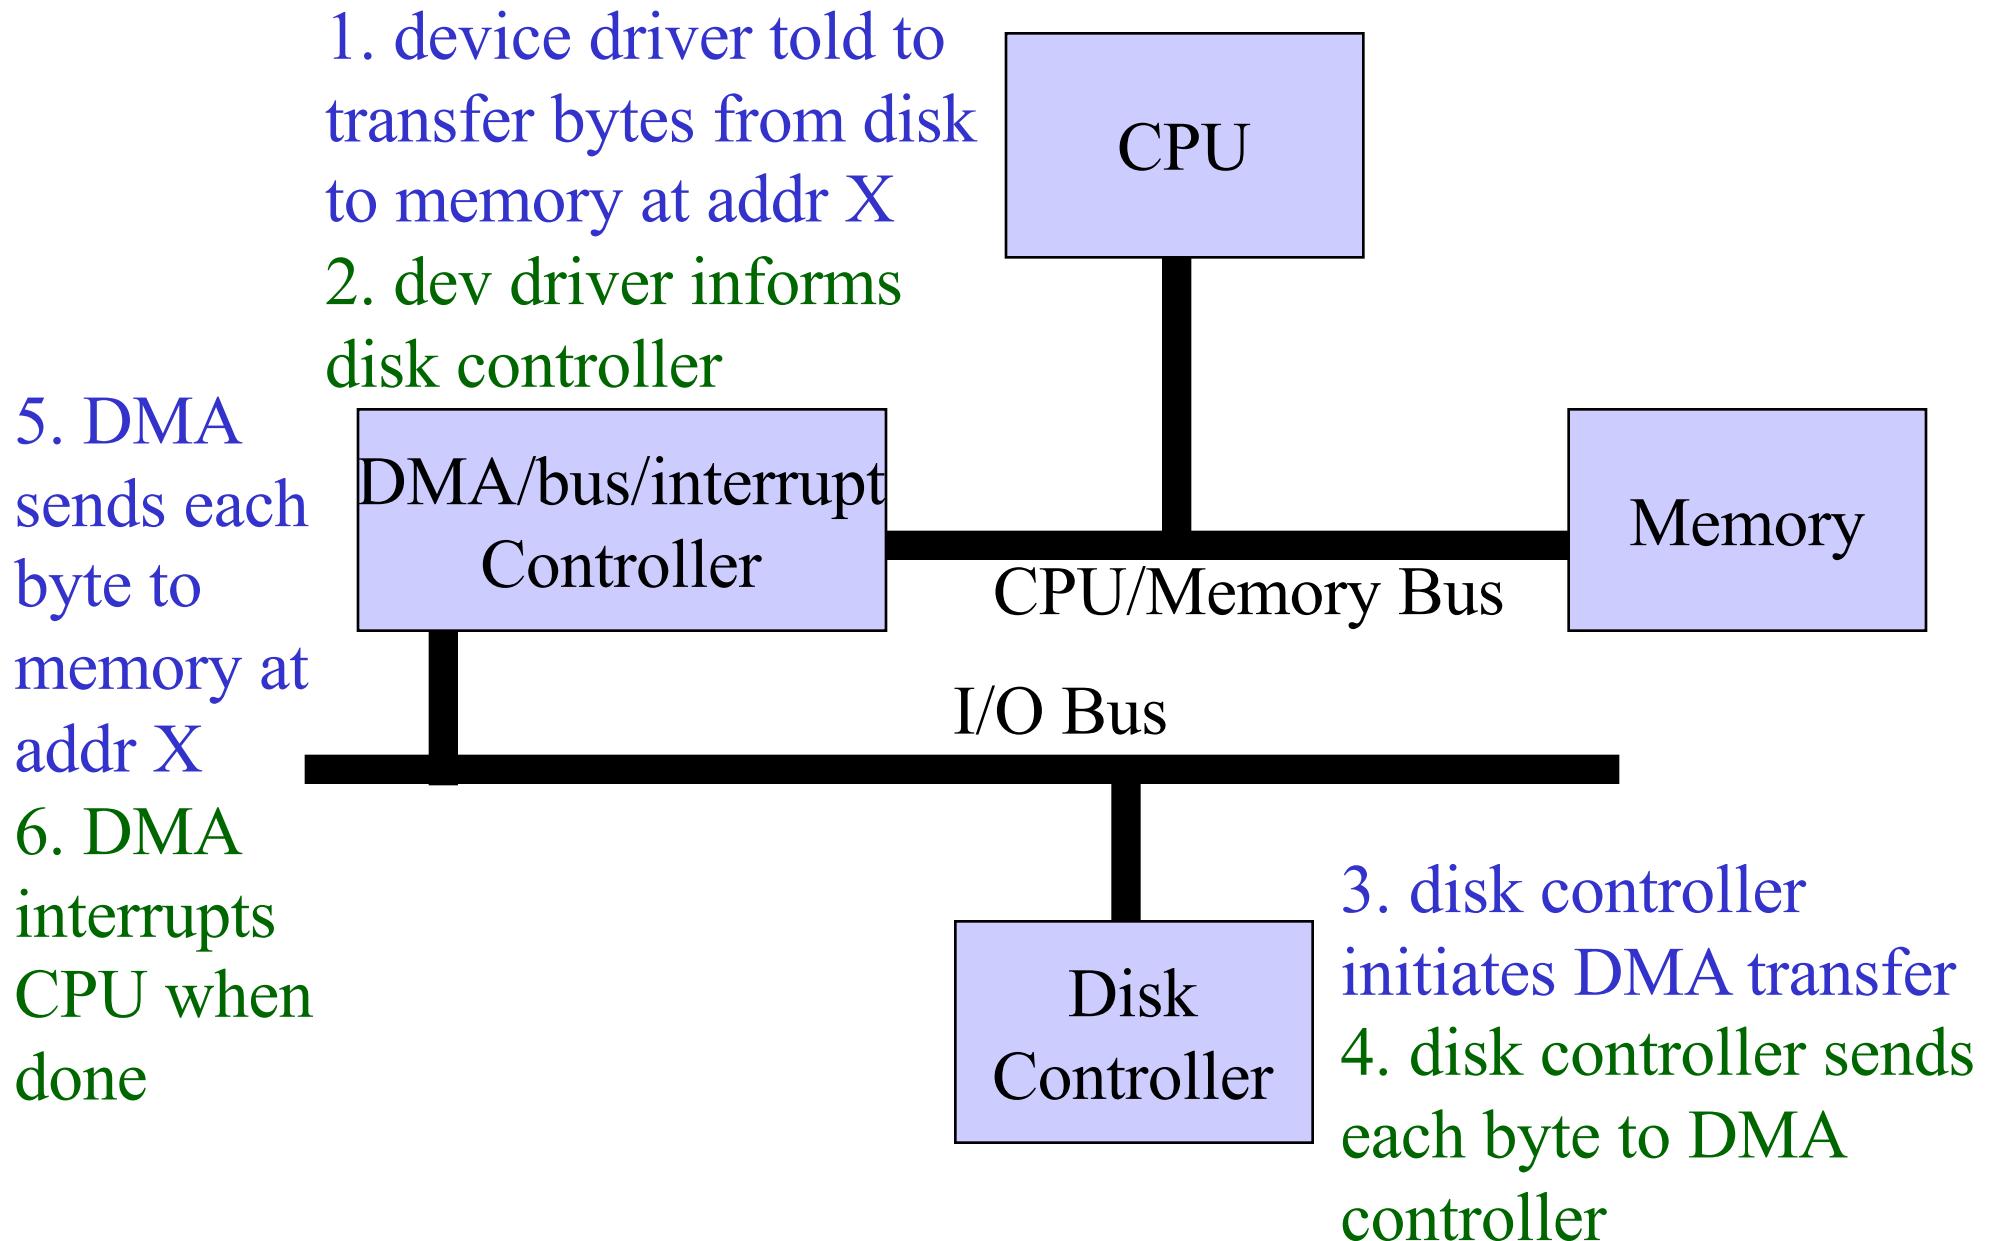
# Interrupt-Driven I/O Operation



# Direct Memory Access (DMA)

- DMA solution: Bypass the CPU for large data copies, and only raise an interrupt at the very end of the data transfer, instead of at every intermediate block
- Modern systems offload some of this work to a special-purpose processor, Direct-Memory-Access (DMA) controller
- The DMA controller operates the memory bus directly, placing addresses on the bus to perform transfers without the help of the main CPU

# DMA with Interrupts Example



# Direct Memory Access (DMA)

- Since both CPU and the DMA controller have to move data to/from main memory, how do they share main memory?
  - Burst mode
    - While DMA is transferring, CPU is blocked from accessing memory
  - Interleaved mode or “cycle stealing”
    - DMA transfers one word to/from memory, then CPU accesses memory, then DMA, then CPU, etc...
      - interleaved
  - Transparent mode – DMA only transfers when CPU is not using the system bus
    - Most efficient but difficult to detect

# Memory-Mapped I/O

- Non-memory mapped (port or port-mapped) I/O typically requires special I/O machine instructions to read/write from/to device controller registers
  - e.g. on Intel x86 CPUs, have IN, OUT
    - Example: OUT dest, src (using Intel syntax, not Gnu syntax)
      - Writes to a device port dest from CPU register src
    - Example: IN dest, src
      - Reads from a device port src to CPU register src
    - Only OS in kernel mode can execute these instructions
    - Later Intel introduced INS, OUTS (for strings), and INSB/INSW/INSD (different word widths), etc.

# Memory-Mapped I/O (2)

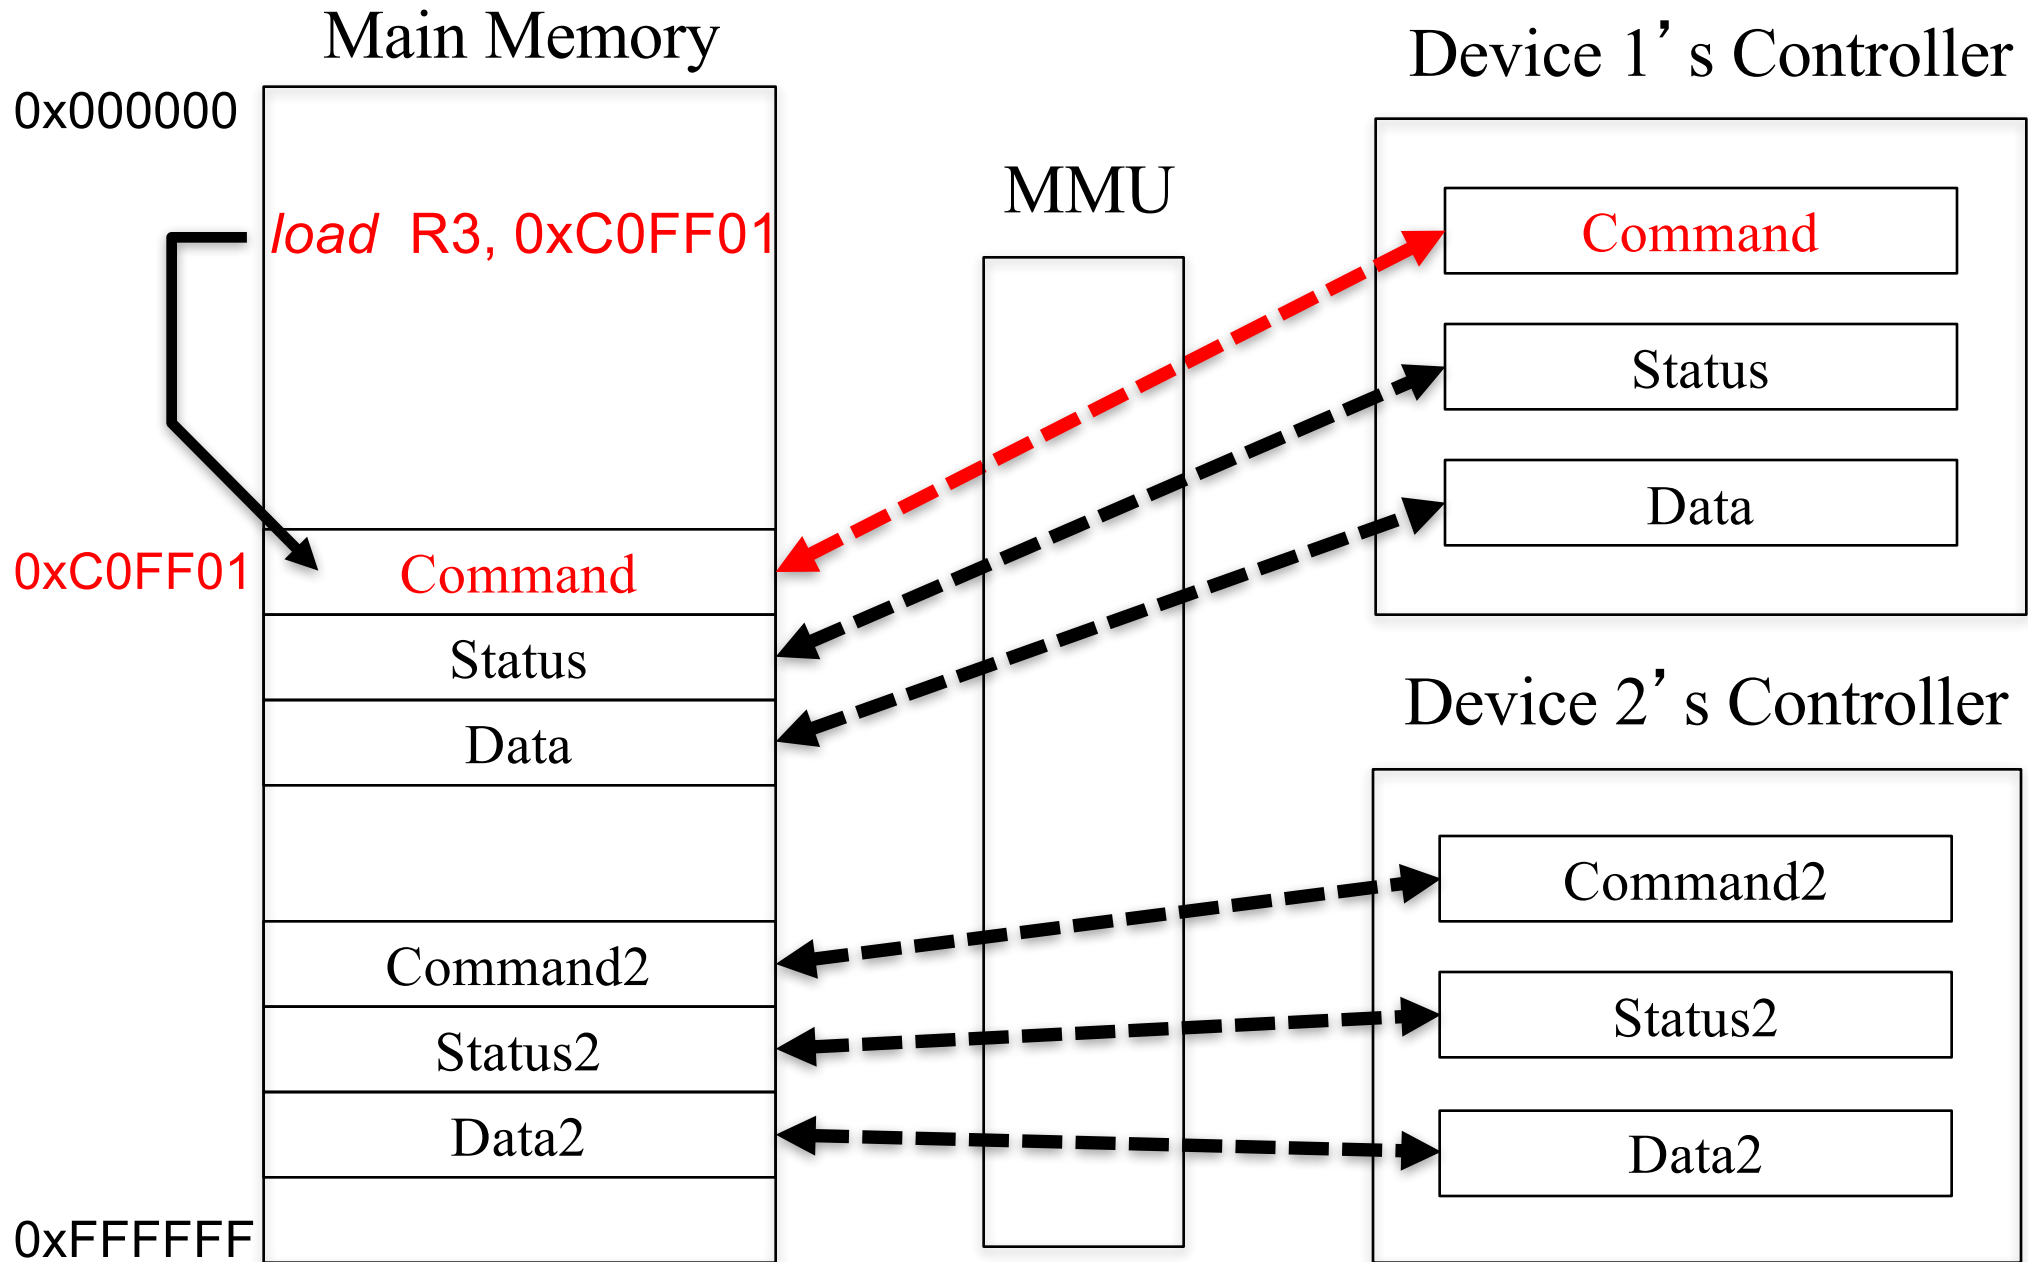
- port-mapped I/O is quite limited
  - IN and OUT can only store and load
  - don't have full range of memory operations for normal CPU instructions
    - Example: to increment the value in say a device's data register, have to copy register value into memory, add one, and copy it back to device register.
    - AMD did not extend the port I/O instructions when defining the x86-64

# Memory-Mapped I/O (3)

- Memory-mapped I/O: device registers and device memory are mapped to the system address space
- With memory-mapped I/O, just address memory directly using normal instructions to speak to an I/O address
  - e.g. `load R3, 0xC0FF01`
  - the memory address `0xC0FF01` is mapped to an I/O device's register
- Memory Management Unit (MMU) maps memory values and data to/from device registers
  - Device registers are assigned to a block of memory
  - When a value is written into that I/O-mapped memory, the device sees the value, loads the appropriate value and executes the appropriate command



# Memory-Mapped I/O (4)



# Memory-Mapped I/O (5)

- Typically, devices are mapped into lower memory
  - frame buffers for displays take the most memory, since most other devices have smaller buffers
  - Even a large display might take only 10 MB of memory, which in modern address spaces of tens-hundreds of GBs is quite modest – so memory-mapped I/O is a small penalty

## Device I/O Port Locations on PCs (partial)

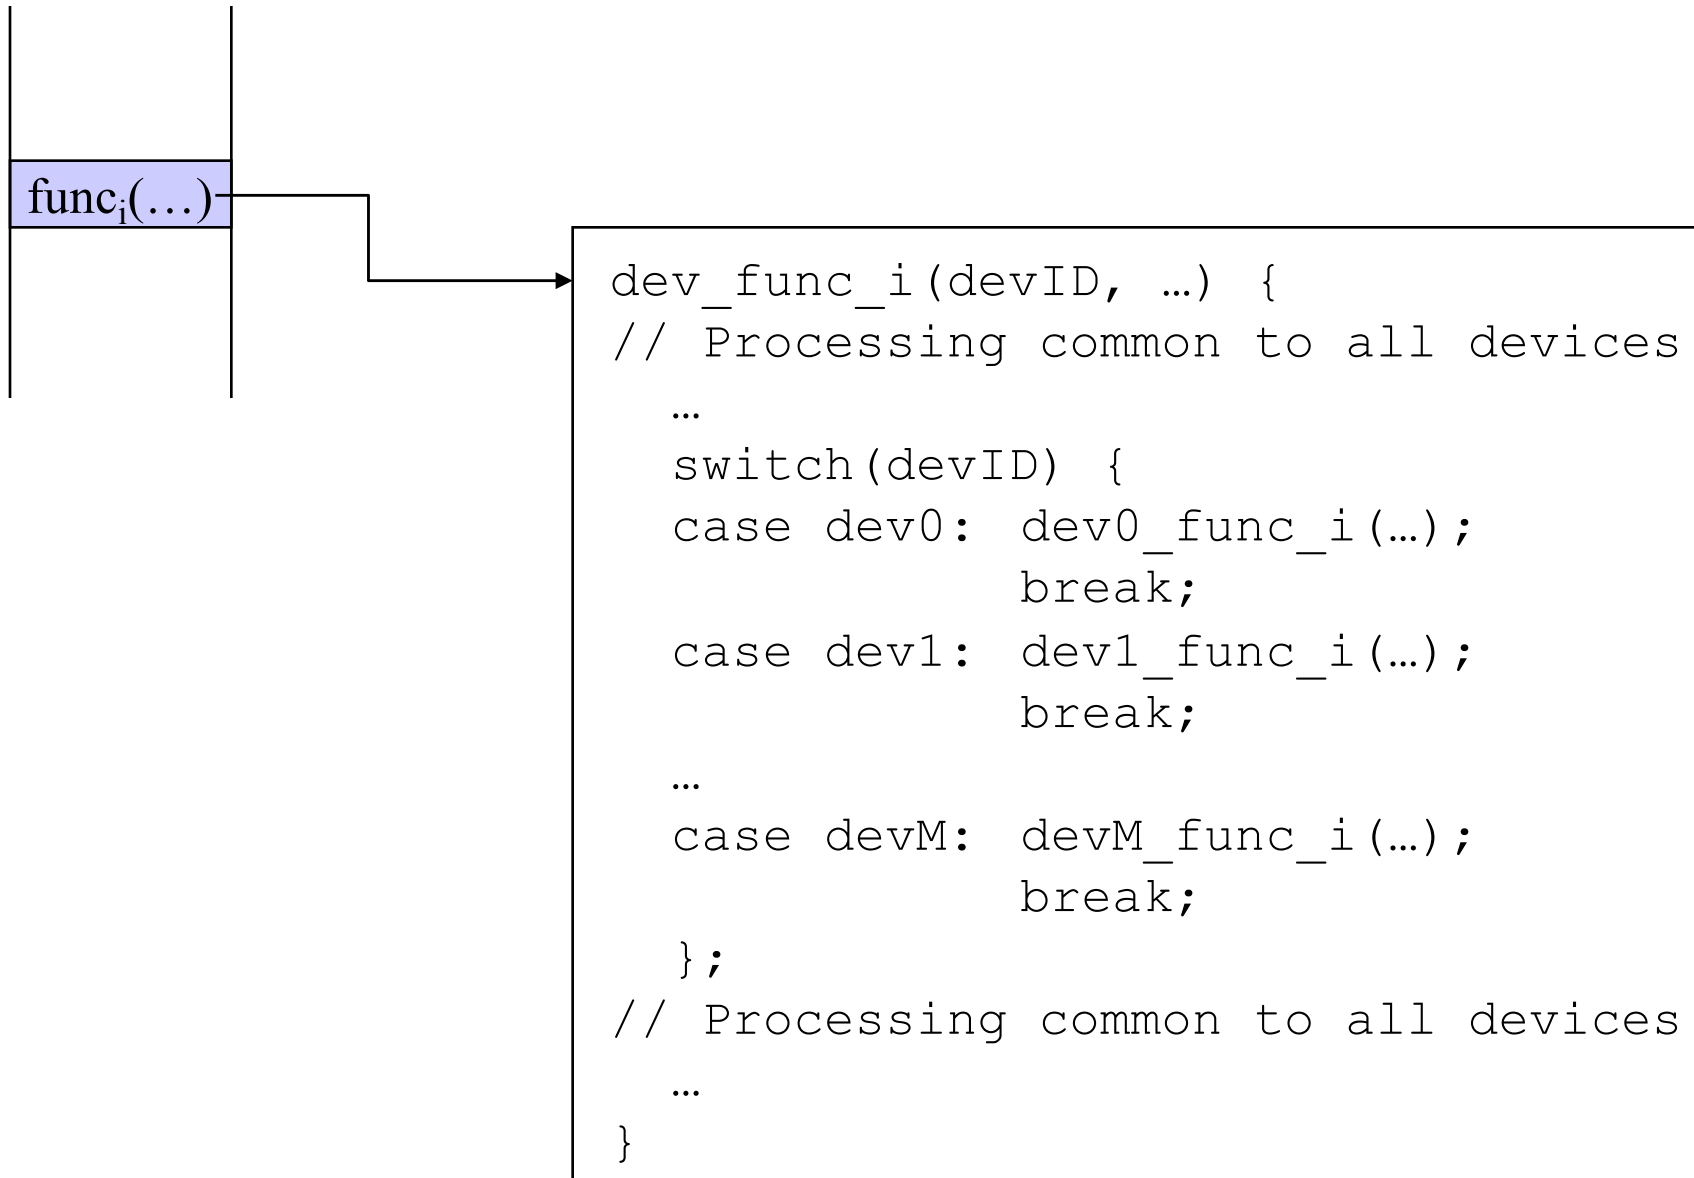
I/O address range (hexadecimal)	device
000–00F	DMA controller
020–021	interrupt controller
040–043	timer
200–20F	game controller
2F8–2FF	serial port (secondary)
320–32F	hard-disk controller
378–37F	parallel port
3D0–3DF	graphics controller
3F0–3F7	diskette-drive controller
3F8–3FF	serial port (primary)

# Device Independent Part

- A set of system calls that an application program can use to invoke I/O operations
- A particular device will respond to only a subset of these system calls
  - A keyboard does not respond to *write()* system call
- POSIX set: *open()*, *close()*, *read()*, *write()*, *lseek()* and *ioctl()*

# Device Independent Function Call

Trap Table



# Adding a New Device

- Write device-specific functions for each I/O system call
- For each system call, add a new *case* clause to the *switch* statement in device independent function call
- Compile the kernel and new drivers

Problem: Need to recompile the kernel, every time a new device or a new driver is added